Attacks on Virtual Machine Emulators

Peter Ferrie, *Senior Principal Researcher, Symantec Advanced Threat Research*
peter_ferrie@symantec.com

Abstract **As virtual machine emulators have become commonplace in the analysis of malicious code, malicious code has started to fight back. This paper describes known attacks against the most widely used virtual machine emulators (*VMware* and *VirtualPC*). This paper also demonstrates newly discovered attacks on other virtual machine emulators (*Bochs*, *Hydra*, *QEMU*, and *Xen*), and describes how to defend against them.**

Index Terms **Hardware-assisted, Hypervisor, Para-virtualization, Virtual Machine**

I. INTRODUCTION

Virtual machine emulators have many uses. For anti-malware researchers, the most common use is to place unknown code inside a virtual environment, and watch how it behaves. Once the analysis is complete, the environment can be destroyed, essentially without risk to the real environment that hosts it. This practice provides a safe way to see if a sample might be malicious.

The simplest attack that malicious code can perform on a virtual machine emulator is to detect it. As more security researchers rely on virtual machine emulators, malicious code samples have appeared that are intentionally sensitive to the presence of virtual machine emulators. Those samples alter their behavior (including refusing to run) if a virtual machine emulator is detected. This behaviour makes analysis more complicated, and possibly highly misleading. Some descriptions and samples of how virtual machine emulators are detected are presented in this paper.

A harsher attack that malicious code can perform against a virtual machine emulator is the denial-of-service; specifically, this type of attack causes the virtual machine emulator to exit. Some descriptions and samples of how that is done are presented in this paper.

Finally, the most interesting attack that malicious code can perform against a virtual machine emulator is to escape from its protected environment. No examples of this type of attack are presented in this paper.

It is important to note here that most virtual machine emulators are not designed to be completely transparent. They are meant to be "good enough" so that typical software can be fooled to run inside them. Their use in the analysis of malicious code was never a requirement. This situation is changing, though, with the creation of new virtual machine emulators, such as *Hydra* [1]. However, even with full knowledge of what has been used to detect existing virtual

machine emulators, it is clearly difficult to develop a virtual machine emulator that cannot be detected. Some descriptions and samples of how to detect *Hydra* are included in this paper.

The interest in detecting virtual machine emulators is also not limited to the authors of malicious code. If malicious code is released that makes use of its own virtual machine emulator, then it will become necessary for anti-malware researchers to find ways to detect the virtual machine emulator, too.

Sample detection code is presented in the appendix. For simplicity and to prohibit trivial copying, only 16-bit real mode assembler code for .COM-format files is supplied.

Virtual machine emulators come in two forms: "hardware-bound" (also known as para-virtualization) and "pure software" (via CPU emulation). The "hardware-bound" category can be split into two subcategories: "hardware-assisted" and "reduced privilege guest" (or ring 1 guest).

Both forms of the hardware-bound virtual machine emulators rely on the real, underlying CPU to execute non-sensitive instructions at native speed. They achieve better performance, for this reason, when compared with pure software implementations. However, since they execute instructions on a real CPU, they must make some changes to the environment, in order to share the hardware resources between the guest operating system and the host operating system. Some of these changes are visible to applications within the guest operating system, if the applications know what those changes look like.

SECTION 1: HARDWARE

II. HARDWARE-BOUND VIRTUAL MACHINE EMULATORS

The difference between hardware-assisted virtual machine emulators and reduced privilege guest virtual machines emulators is the presence of virtual machine-specific instructions in the CPU. The hardware-assisted virtual machine emulators use CPU-specific instructions to place the system into a virtual mode. The guest runs at the same privilege level that it would do if it truly controlled the CPU in the absence of the virtual machine emulator. The important data structures and registers have shadow copies that the guest sees, but these shadow copies have no effect on the host.

Instead, the host controls the real data structures and registers. The result is that the virtualization is almost completely transparent. The host can direct the CPU to notify

it of specific events, such as an attempt to query the capabilities of the underlying CPU, or to access particular memory locations and important registers.

By contrast, the reduced privilege guest virtual machine emulators must virtualize the important data structures and registers themselves. The guest is run at a lower privilege level than it would do if it truly controlled the CPU. There is no way to prevent the CPU from notifying the host of all interesting events.

The idea of hardware-bound virtual machine emulators is not new - *IBM* has been using them for four decades on the *System/360* hardware and its descendants.

In the days of DOS, reduced privilege guest virtual machine emulators could be implemented by hooking interrupt 1 [2], for example. The interrupt 1 hook allows the real CPU to execute instructions at native speed, but the downside is that every instruction is also treated as though it were sensitive.

Another method of reduced-privilege guest virtual machine emulation is buffered code emulation [3]. Buffered code emulation works by copying an instruction into a host-controlled buffer and executing it there, if it is not a sensitive or special instruction. Buffered code emulation has fairly good performance.

A major problem for both of these methods when implemented in DOS, is that DOS has no notion of privileges. Thus, reduced privilege guest is actually a misnomer since it runs at the same privilege level as the host. As a result, code could "escape" from the environment by hooking an "Interrupt ReQuest Vector" (IRQ) and then waiting for that IRQ to be asserted (or, in the case of disk drive IRQs, issuing a command which causes the IRQ to be asserted on completion). There were also problems when the emulation was run in virtual-8086 mode, because the emulator couldn't switch into protected mode and retain control.

This is not a problem for more modern operating systems, though, such as *Windows* and *Linux*. In fact, *VirtualPC* [4] uses buffered code emulation. It preloads up to 128 bytes, and executes them from there, if possible. Otherwise, it wraps special code around them, and then it passes them to the *VMM.sys* driver that performs the actual execution. The use of buffered code emulation allows *VirtualPC* to intercept instructions that cannot be intercepted by other hardware-bound virtual machine emulators.

Another application that uses buffered code emulation is *Dynamo Rio* [5]. The difference between *VirtualPC* and *Dynamo Rio* in this case is that *Dynamo Rio* runs at an application level and as a Dynamic Link Library within the process space of the guest application, whereas *VirtualPC* runs at the system level. *Dynamo Rio* actively attempts to hide itself by intercepting and manipulating memory requests, module lists, etc. Since it is not a virtual machine emulator as

defined by the terms described in the introduction, it was not considered further.

Some examples of reduced privilege guest virtual machine emulators are *VMware* [6], *Xen* [7], and *Parallels* [8]. One other product called *Virtuozzo* [10] is known to the author, but a copy could not be acquired at the time of writing. It is likely to be a reduced privilege guest virtual machine emulator.

III. HARDWARE-ASSISTED VIRTUAL MACHINE EMULATORS

Xen 3.x, *Virtual Server 2005* [9], and *Parallels*, can exist as hardware-assisted virtual machine emulators. Virtuozzo might also exist as a hardware-assisted virtual machine emulator.

From a malicious code author's perspective, the most interesting thing about hardware-assisted virtual machine emulators (hypervisors) is that they can be used to virtualize the currently running operating system at any point in time. Thus, the host can boot to completion, and launch any number of applications as usual, with one them being the virtual machine emulator. That emulator then sets up some CPU-specific control structures and uses the *VMLAUNCH* (*Intel*) or *VMRUN* (*AMD*) instruction to place the operating system into a virtualized state. At that point, there are effectively two copies of the operating system in existence, but one (the host) is suspended while the other (the guest) runs freely in the new state. Whenever an interesting event (an intercept, interrupt, or exception) occurs, the host operating system (the virtual machine emulator) regains control, handles the event, and then resumes execution of the guest operating system.

Thus, any machine that supports the existence of a hypervisor can have a hypervisor start running at any time. Neither the operating system, nor the user, will be aware of it. Further, the hypervisor is actually more privileged than the operating system itself, since it sees the interesting events first and can hide them even from the host operating system. A hypervisor is, in effect, an "enhanced privilege host". Additionally, once a hypervisor is active, no other hypervisor installed later can gain full control of the system. The first hypervisor is in ultimate control.

In theory, once the guest is active, the virtual machine emulator cannot be detected since it can intercept all sensitive instructions, including the *CPUID* instruction. The instructions that would leak information now see a shadow copy of the sensitive information which appears to correspond to a real CPU. The suggested methods to hide the presence of the hypervisor are: clear the *CPUID* flag that corresponds to the hardware-assisted "Virtual Machine eXtensions" (VMX) capabilities or emulate the VMX instructions, which would allow for nested virtual machines. The former method is apparently used by *BluePill*; the latter method is used by *Xen*.

The method used by *Xen* is especially interesting since it means that even a hypervisor can be fooled into thinking that

it is running on the real hardware. Normally, one might think that if a hypervisor starts running correctly, then it is in full control of the system. In fact that is not the case.

This promise of "undetectibility" has alarmed many people. Early *Intel* documentation regarding these Virtual Machine Extensions went as far as to say that it was impossible to detect. More recent documentation has softened the language to say that it is difficult to detect. It is indeed difficult to detect, but not impossible.

The most obvious attack against hypervisors is to check a local time source, such as the "Time Stamp Counter" (TSC). This fact was understood by both *Intel* and *AMD*. The result is the "TSCDelta" field in the "Virtual Machine Control Block" (VMCB) which can be used to skew the guest's TSC by an appropriate value to hide the delay caused by faults to the hypervisor.

Therefore, all of the currently documented methods for detecting hypervisors rely on external timing. Specifically, they rely on the fact that executing certain instructions many times will take far longer within a hypervisor environment than without [11]. While that is true, without any baseline comparison (time required for the same machine to run the same number of iterations of the same instructions, prior to the hypervisor being installed), it is impossible to know that a hypervisor is present. Any other time source must be considered suspect. For example, the protocol for interacting with time servers is documented and easily intercepted by the hypervisor.

An alternative exists that relies on a different kind of timing. The method was discovered earlier this year, but no details were given at that time [12]. The method is described below for the first time.

The "Translation Lookaside Buffers" (TLBs) can be filled with known data, by accessing a series of present pages. Then if a hypervisor is present, a hypervisor event can be forced to occur by using a hypervisor-sensitive instruction.

In particular, we need a hypervisor-sensitive instruction that is not otherwise destructive to the TLBs. There is only one instruction that meets the criteria: *CPUID*. *CPUID* is the only instruction that is intercepted by a hypervisor, is not privileged, and most importantly, does not affect memory in any way.

If the TLBs are explicitly flushed, then the time to access a new page can be determined by reading the time stamp counter before and after the access. This duration can be averaged over the number of TLBs to be filled. Once the TLBs are filled, the time to access a cached page can be determined by reading the time stamp counter before and after the access of each page in the TLBs. This duration can also be averaged over the number of TLBs that were filled.

Next, the *CPUID* instruction is executed, which will cause a hypervisor intercept to occur, and at least some of the TLBs will be flushed as a side-effect. If a hypervisor event occurred, then each of the pages that should be in the TLBs can be accessed again, and the access time can be measured. If the access time matches that of a new page instead of a cached page, then the hypervisor's presence is revealed.

IV.   PURE SOFTWARE VIRTUAL MACHINE EMULATORS

Pure software virtual machine emulators work by performing equivalent operations in software for any given CPU instruction. The main advantage that pure software virtual machine emulators have over hardware-bound virtual machines is that the pure software CPU does not have to match the underlying CPU. This allows a guest environment to be moved freely between machines of different architectures. Some examples of pure software virtual machine emulators are *Hydra*, *Bochs* [13], and *QEMU* [14].

Another method of virtual machine emulation is most often used by anti-virus software. It emulates both the CPU and a portion of an operating system, such as *Windows* or *Linux*. Two examples of this are *Atlantis* [15] and *Sandbox* [16]. Both of these are intended to allow a malicious file to "run", while capturing information about its behavior in a completely safe manner. *Atlantis* supports DOS, *Windows*, and *Linux*. *Sandbox* supports *Windows* only.

Some virtual machine emulators, such as *Hydra*, *Bochs, and Atlantis*, support different CPUs internally, in order to more reliably emulate an environment when the required CPU is not known. A problem for any emulator is that different generations of CPUs can display slightly different behaviors for identical instructions. For *Intel* 80x86 CPUs, for example, the *AAA* instruction sets the flags in one of three different ways, depending on whether the CPU is an 80486 or *Pentium*, a *Pentium 2* or *Pentium 3*, or a *Pentium 4* or later. Therefore, if a pure software virtual machine emulator is written for one specific CPU, the software that is emulated might not behave correctly. This is, of course, also a problem for hardware-bound virtual machine emulators, but more so in their case because they cannot do anything about it.

V.   VIRTUAL MALICIOUS CODE

Predictably, the increasing interest in virtualization has led some researchers to propose malicious uses for virtual machines. One reduced privilege guest virtual machine rootkit, called *SubVirt*, has been described in detail elsewhere [17], and is described briefly here. *SubVirt* works by installing a second operating system. This operating system becomes the new host operating system, which carries an operating system-specific virtual machine emulator. *SubVirt* supports both the *Windows* and *Linux* operating systems. For the *Windows* platform, *SubVirt* carries *VirtualPC*; for the *Linux* platform, *SubVirt* carries *VMware*. Once the new host operating system loads and runs the virtual machine emulator,

the virtual machine emulator places the old host operating system into a virtual machine and carries on as before. In the absence of software that is able to recognize the presence of a virtual machine emulator, software within the system will not easily determine that the system has been compromised.

Two hardware-assisted virtual machine rootkits have also been described elsewhere, by their authors. One is *BluePill* [18], and the other is *Vitriol* [19]. Both of them work by making use of the virtual machine extensions that exist in newer *AMD* and *Intel* CPUs respectively.

It seems that none of these applications is available to other anti-malware researchers.

VI. DETECTING *VMWARE*

VMware is a proprietary, closed-source, reduced privilege guest virtual machine emulator. It supports guest-to-host and host-to-guest communication. Since it relies on the underlying hardware for execution of instructions, it must relocate sensitive data structures, such as the Interrupt Descriptor Table and the Global Descriptor Table. *VMware* also makes use of the Local Descriptor Table which is not otherwise used by *Windows*. Thus, a simple detection method for *VMware* is to check for a non-zero Local Descriptor Table (LDT) base on *Windows* [20]. The more common method for detecting *VMware* is to check the value of the Interrupt Descriptor Table, using the "RedPill" [21] method. For the "RedPill" method, if the value of the Interrupt Descriptor Table base exceeds a certain value, a virtual machine emulator is assumed to be present. However, as the LDT paper shows, this method is unreliable on machines with multiple CPUs. The "Scooby Doo" [22] method uses the same basic idea as the RedPill method but it compares the Interrupt Descriptor Table base value to specific hard-coded values in order to identify *VMware* specifically. While the Scooby Doo method is less likely to trigger false positives, compared to the RedPill method, there is still the chance that some false positives will occur.

In addition to the Descriptor Table methods, *VMware* offers a method of guest-to-host and host-to-guest communication which can also be used to detect the presence of *VMware*. The most common form of this detection is the following [23]:

```
mov eax, 564d5868h  ;'VMXh'
mov ecx, 0ah        ;get VMware version
mov dx, 5658h       ;'VX'
in  eax, dx
cmp ebx, 564d5868h  ;'VMXh'
je  detected
```

When run in ring3 of a protected-mode operating system, such as *Windows* or *Linux*, execution of the *IN* instruction causes an exception to be generated, unless the I/O privilege level is altered. This is because the *IN* instruction is a privileged instruction. The reason that the Interrupt Descriptor Table is relocated is to hook this exception privately. The exception can be normally trapped by an application. However, if *VMware* is running, no exception is generated. Instead, the EBX register is altered to contain 'VMXh' (the ECX register is also altered to contain the *VMware* product ID, which is not relevant in this case).

This detection method was attempted recently in the W32/Polip virus [24]. The virus author attempted to obfuscate it and ended up by introducing a bug, so *VMware* was not detected even when it was running.

Of course, other values in the ECX register can be specified for different effects [25]. Since the execution of the *IN* instruction should never change register values other than the EAX register in a real machine, disabling the "get *VMware* version" method alone will not be sufficient to hide *VMware*.

There are many other ways to detect the presence of *VMware*, depending on the guest operating system that is in use. For example, the *Windows* registry is full of *VMware*-specific keys, but all of these can be removed. Other methods depend on the presence of particular hardware, such as hard disks whose device names are constant, and network cards whose MAC addresses fall within a predictable range. The problem with these dependencies is that, depending on the intended use of the virtual system, none of these hardware elements might be present, and some of them require special privileges to access.

Going beyond detection, in December 2005, it was disclosed that a component of *VMware* allowed an attacker to escape from the environment. Specifically, the "VMnat" contained an unchecked copy operation while processing specially crafted 'EPRT' and 'PORT' FTP requests [26]. The result was heap buffer corruption within the host environment, with the potential to execute arbitrary code there.

A more serious vulnerability potentially exists in hardware-bound virtual machine emulators, if the guest can interact with third-party devices on the system. For example, if a buffer-overflow vulnerability exists in a network driver in the host environment, it might be possible for an application within the guest environment to send a specially crafted network packet that reaches the host network driver intact, and thus exploit that vulnerability.

VII. DETECTING *VIRTUALPC*

VirtualPC is a proprietary, closed-source, reduced privilege guest virtual machine emulator. It supports guest-to-host and host-to-guest communication. A version exists for the *Macintosh* platform, as well as for the *Windows* platform. Only the *Windows* version is considered here.

Just like *VMware*, *VirtualPC* must relocate sensitive data structures, such as the Interrupt Descriptor Table and the Global Descriptor Table. Just like *VMWare*, VirtualPC makes use of the Local Descriptor Table. Thus, RedPill, LDT, and Scooby Doo, all work to detect VirtualPC.

Whereas *VMware* uses a special port to perform guest-to-host and host-to-guest communication, *VirtualPC* relies on the execution of illegal opcodes to raise exceptions that the kernel will catch. This method is very similar to the illegal opcode execution that *Windows NT* and later operating systems use in their DOS box to communicate with the operating system. By reverse-engineering the *VirtualPC* executable file, the author found that the opcodes are the following:

```
0F 3F x1 x2
0F C7 C8 y1 y2
```

In ordinary circumstances, execution of these opcodes causes an exception to be generated. The 0F 3F opcode causes an exception because it is an otherwise undefined opcode. The 0F C7 C8 opcode causes an exception because it is an illegal encoding of an existing opcode. This exception can be trapped by an application. However, if *VirtualPC* is running, no exception is generated, depending on the values of x1, x2, y1, and y2.

The full list of allowed values for x1 and x2 is not known. However, the BIOS code in *VirtualPC* uses the values 0A 00, 11 00, 11 01, and 11 02. The file-sharing module that can be installed uses value 02 followed by 01-13, and 07 0b. These appear to be examples of guest-to-host communication. An example of host-to-guest communication is given in the following: if x1 is 03 and x2 is 00, then the current host time (in hour:minute:second notation) is placed into the DX, CX, and AX, registers respectively (see *VIRTUALPC* TIME demo). Other values for x1 and x2, such as 02 00, return other values in the CPU registers. The values 10 01-03 and 10 06 alter the Z flag. The IsRunningInsideVirtualMachine() API uses the value 07 0B.

The allowed values for y1 are 00-04. The allowed values of y2 depend on the value of y1. If y1 is 00 or 03, then y2 can be 00-03. If y1 is 01, then y2 can be 00-02. If y1 is 02, then y2 can be 00-04. If y1 is 00, then y2 can only be 00. The BIOS code in *VirtualPC* uses the values 00 00 and 00 01. The Virtual Machine Additions driver uses the value 00 01. The IsRunningInsideVirtualMachine() API uses the value 01 00.

Another method for detecting *VirtualPC* relies on the fact that *VirtualPC* does not limit the length of an instruction. *Intel* and *AMD* CPUs have a maximum instruction length of 15 bytes. This is achievable only in 16-bit mode, using the 81 opcode. The instruction would look something like the following:

```
lock
add dword ptr cs:[eax+ebx+01234567], 89abcdef
```

In addition to the "ADD" instruction, this encoding of the 81 opcode also supports "OR", "ADC", "SBB", "AND", "SUB", or "XOR". The 81 opcode also supports the "CMP" instruction, but it is not permitted in this context because of the "LOCK" prefix.

Any instruction longer than 15 bytes - which is achievable only by the addition of redundant prefixes - will cause a General Protection Fault. However, *VirtualPC* does not issue this exception, seemingly no matter how long the instruction (see *VIRTUALPC* ILEN demo).

As noted above, *VirtualPC*'s use of buffered code emulation allows it to intercept instructions that cannot be intercepted by other hardware-bound virtual machine emulators, particularly the hardware-based ones. In theory, the RedPill method could be defeated by intercepting the *SIDT* instruction, as described in the *SubVirt* paper. However, this is currently not implemented. The *CPUID* instruction is one instruction that *VirtualPC* does intercept. On a real CPU, the returned vendor identification string is either "GenuineIntel" or "AuthenticAMD". In *VirtualPC*, though, it is "ConnectixCPU", a reference to the company which developed the earlier versions of *VirtualPC*.

As with *VMware*, there are many other ways to detect the presence of *VirtualPC*, including the use of hardware devices with constant names. One detection method is even described by a *Microsoft VirtualPC* developer [28]. That method queries the name of the manufacturer of the motherboard, which is "Microsoft Corporation" in *VirtualPC*. Since there can be only one motherboard, the code can be shortened significantly (see *VIRTUALPC* BOARD demo). However, the problem with this method is that it requires that the *Windows Management Instrumentation* service is running.

VIII.   DETECTING *PARALLELS*

Parallels is a proprietary, closed-source, reduced privilege guest virtual machine emulator. It supports guest-to-host and host-to-guest communication. It resembles *VirtualPC* in many ways. Just like *VirtualPC*, a version exists for the *Macintosh* platform, as well as for the *Windows* platform. Only the *Windows* version is considered here.

Just like *VMware* and VirtualPC, Parallels must relocate sensitive data structures, such as the Interrupt Descriptor Table and the Global Descriptor Table. Just like *VMWare* and *VirtualPC*, *Parallels* also makes use of the Local Descriptor Table. Thus, RedPill, LDT, and Scooby Doo, all work to detect *Parallels*.

Parallels has two methods of guest-to-host and host-to-guest communication. One of them relies on the execution of an opcode to raise an exception. In this case, the opcode is the *BOUND* instruction. The difference between the method used by *Parallels*, and the method used by other virtual machine emulators, is that *Parallels* uses authentication to determine whether or not the exception is trapped by the kernel.

The method of authentication is to pass in the CPU registers (EAX, ECX, EDX, EBX) values that are specific to the currently executing session. When *Parallels* first loads the kernel driver, the driver halts the CPU and waits for an

interrupt to occur. At that time, the *RDTSC* instruction is read sixteen times in a row, and the lowest byte is stored in an array that corresponds to those registers. To communicate with the kernel, the guest sets the EBP registers to the string "0x90", and the EDI register contains the index of the function to execute in a function pointer array, and then executes the *BOUND* instruction with values that are guaranteed to raise the *BOUND* exception. The main *Parallels* executable file also uses this method.

```
pushad
mov     esi, [ebp+xxxx]
mov     eax, [esi] ;load auth value
mov     ebx, [esi+4] ;load auth value
mov     ecx, [esi+8] ;load auth value
mov     edx, [esi+0Ch] ;load auth value
mov     edi, [esi+10h] ;load auth value
mov     esi, [ebp+xxxx] ;load real esi
xor     ebp, ebp
push    ebp ;upper bound value
push    ebp ;lower bound value
mov     ebp, '0x90'
bound   ebp, [esp] ;raise exception
add     esp, 8 ;discard bound values
popad
```

The second method of guest-to-host and host-to-guest communication occurs through the use of the *INT 1B* vector. In that case, the registers are initialized in the following way: the ESI register contains the string "magi", the EDI register contains the string "c!nu", and the EBX register contains the string "mber". It spells "magic!number". The EDX register is set to point to any variables on the stack that must be passed, and the EAX register is set to the function number to call. One of the *Parallels* driver files also uses this method.

```
mov     esi, 'magi'
mov     edi, 'c!nu'
mov     ebx, 'mber'
push    [ebp+xxxx]
push    [ebp+xxxx]
push    [ebp+xxxx]
push    xxxxxxxx
mov     edx, esp
mov     eax, 0
int     1bh
```

The reason for the two different methods is that the *BOUND* method is available from user mode, so it must be protected from abuse by non-privileged applications. The *INT 1B* method is available only from kernel mode, so a user with sufficient privileges to install a kernel-mode driver should presumably have sufficient privileges to communicate with *Parallels* itself.

In addition, the author found not another way to detect *Parallels*, but a way to crash it. By entering v86 mode (a *Windows* DOS box was used) and issuing a *SIDT* instruction with the Trap flag set, *Parallels* encounters a fatal error and closes.[1]

SECTION 2: SOFTWARE

---

[1] The vendor was notified, but did not respond after sixty days.

Pure software virtual machine emulators are also vulnerable to detection. In their case, detection is possible mostly because of software bugs or incomplete support for the CPU which is being emulated.

IX. DETECTING *BOCHS*[2]

Bochs is an Open Source, pure software virtual machine emulator. It does not support guest-to-host or host-to-guest communication since it is intended to behave like a stand-alone machine. It is vulnerable to a number of detection methods. The simplest of these involves the device support. For example, *Bochs* cannot handle floppy disks of non-standard sizes. Attempting to format a 3.5" floppy disk with more than 18 sectors per track, or with sectors other than 512 bytes in size, will cause a kernel panic. As with *VMware* and *VirtualPC*, *Bochs* has constant names for its hardware devices, but again, the presence of these devices cannot be relied upon. Thus, we are left with the CPU as the target for detection. The author discovered a number of methods to detect *Bochs*. Here are some of them:

- *INVD* and *WBINVD* instructions always flush TLBs (see *BOCHS WBINVD* demo)

  The code works by entering paging mode, and then accessing a page. This causes the CPU to place the page's physical address into one of the Translation Lookaside Buffers. When an *INVD* or *WBINVD* instruction is executed inside *Bochs*, the Translation Lookaside Buffers are flushed. Hence, if the same page is marked "not present" then accessed again, a Page Fault occurs. By registering a Page Fault handler prior to executing the *INVD* or *WBINVD* instruction, *Bochs* can be detected. This detection method is available due to what appears to be an oversight.

- *CMPS* instruction flags are not retained while *REP* continues in single-step mode (see *BOCHS CMPS* demo)

- *SCAS* instruction flags are not retained while *REP* continues in single-step mode (see *BOCHS SCAS* demo)

  These two codes begin by setting the carry flag. Then, in the case of the *CMPS* instruction, two ranges of bytes that are known to be identical are compared (the source and destination registers are set to the same value). In the case of the *SCAS* instruction, a single byte, whose value is known to match the destination, is compared to the destination. The source register is set to the value in memory that is pointed to by the destination register. In a real

---

[2] This list is the longest in this paper because *Bochs* was the first application to be examined, and received the most scrutiny. It does not reflect the quality of the software.

machine, the carry flag remains set until the *REP* has completed. However, in *Bochs*, the flag is updated immediately. By registering a trap handler prior to executing the *CMPS* or SCAS instruction, the carry flag can be seen to have been cleared, and thus *Bochs* can be detected. This detection method is available due to what appears to be an oversight.

- *CPUID* instruction returns wrong value for processor name on *AMD* CPU (see *BOCHS CPUID_AMD1* demo)

  This code works by executing the *CPUID* instruction to check for an *AMD* CPU. If one is found, then the *CPUID* instruction is executed again to query maximum input value for the extended *CPUID* information. If the processor brand string is supported, then the *CPUID* instruction is executed again to query the processor brand string. For a real *AMD K7* processor (the only one that *Bochs* supports), the returned string is "AMD Athlon(tm) P[rocessor]". For *Bochs*, it is "AMD Athlon(tm) p[rocessor]" (note the lowercase 'p'). This detection method is available due to what appears to be an oversight.

- *CPUID* instruction returns wrong value for Easter egg on *AMD* CPU (see *BOCHS* and *QEMU CPUID_AMD2* demo)

  This code works by executing CPUID to check for an *AMD* CPU. If one is found, then the *CPUID* instruction is executed again to query the Easter egg. For a real *AMD K7* processor (the only one that *Bochs* supports), the returned value is "IT'S HAMMER TIME". For *Bochs*, nothing is returned. This detection method is available due to what appears to be an oversight.

- *ARPL* instruction destroys upper 16 bits of 32-bit register in 32-bit mode (see *BOCHS ARPL* demo)

  This code executes the *ARPL* instruction using the undocumented 32-bit register mode. Officially, the instruction accepts 16-bit registers. For some reason, *Bochs* ORs the top 16 bits with 0ff3f0000h, but the author found no real CPU where that behavior occurs. This detection method is available due to what appears to be an oversight.

- 16-bit segment wraparound is not supported (see *BOCHS* and *HYDRA* SEGLOAD demo)

  This code executes a segment:register load, at an offset where the register part is at a lower address than is the segment part. By registering a trap handler prior to executing the load instruction, an exception will occur in *Bochs* that should not occur

at all. Thus *Bochs* can be detected. This detection method is available due to what appears to be an oversight.

- Non-ring0 SYSENTER CS MSR causes kernel panic

  This is similar to the v86 *SIDT* problem in *Parallels*, in that it is not a method to detect *Bochs*, but a way to crash it. By simply writing to the SYSENTER CS MSR (174h) a value with any of the low two bits set, *Bochs* will encounter a kernel panic and close. A real CPU will accept this value since no checks are done until the *SYSENTER* instruction is actually executed. This detection method is available due to what appears to be an oversight.

X.   DETECTING *HYDRA*[3]

Hydra is a proprietary, closed-source, pure software virtual machine emulator. It supports guest-to-host communication, even though it is intended to behave like a stand-alone machine. It does not intentionally support host-to-guest communication. The guest-to-host communication channel exists for the use of plug-ins that can alter the environment and control the execution flow. However a plug-in is not supposed to communicate with the guest. *Hydra* also uses a special port for guest-to-host communication, much like *VMware* does. The key differences between *VMware* and *Hydra* are that in *Hydra*, the port to use is specific to the plug-in; and a plug-in can still cause an exception to be generated, thus better hiding the interaction. Since no host-to-guest communication occurs, no *Hydra*-specific information is returned by the port access. In any case, the author discovered a number of methods to detect *Hydra*. Some of the methods are described in the following:

- *REP MOVS* instruction integer overflow (see *HYDRA MOVS* demo)

- *REP STOS* instruction integer overflow (see *HYDRA STOS* demo)

  This code works by causing a loop counter to overflow, when converting from a dword count to a byte count. Thus no bytes are copied (in the case of the MOVS instruction) or stored (in the case of the *STOS* instruction). This leads the emulator to believe that an error occurred, so a General Protection Fault is raised. In the absence of a General Protection Fault handler, a Double Fault occurs. In the absence of a Double Fault handler, a Triple Fault occurs, leading to the emulator exiting completely. This detection method is available due to a limitation in the string acceleration code.

---

[3] All of the the problems described here have since been fixed.

- 16-bit segment wraparound is not supported (see *BOCHS* and *HYDRA* SEGWRAP demo)

  This code executes a segment:register load, at an offset where the register part is at a lower address than is the segment part. By registering a trap handler prior to executing the load instruction, an exception will occur in *Hydra* that should not occur at all, and thus *Hydra* can be detected. This detection method is available due to what appears to be an oversight.

XI. DETECTING *QEMU*

QEMU is an Open Source, pure software virtual machine emulator. It does not support guest-to-host or host-to-guest communication since it is intended to behave like a stand-alone machine. It supports dynamic translation of code to improve the performance on the supported CPUs. The use of dynamic translation is always risky in the presence of self-modifying code, especially when non-intuitive CPU behavior occurs, such as a self-overwriting *REP* sequence[4]. The author discovered a number of methods to detect QEMU. Some of the methods are described in the following:

- *CPUID* instruction returns wrong value for processor name on *AMD* CPU (see *QEMU CPUID_AMD* demo)

  This code works by executing the *CPUID* instruction to check for an *AMD* CPU. If one is found, then the *CPUID* instruction is executed again to query maximum input value for the extended *CPUID* information. If the processor brand string is supported, then the *CPUID* instruction is executed again to query the processor brand string. For a real *AMD K7* processor, the returned string is "AMD [processor name] Processor". For *QEMU*, it is "QEMU Virtual CPU version x..x..x".

- *CPUID* instruction returns wrong value for Easter egg on *AMD* CPU (see *BOCHS* and *QEMU CPUID_AMD2* demo)

  This code works by executing the *CPUID* instruction to check for an *AMD* CPU. If one is found, then the *CPUID* instruction is executed again

to query the Easter egg. For a real *AMD K7* processor, the returned value is "IT'S HAMMER TIME". For *QEMU*, nothing is returned. This detection method is available due to what appears to be an oversight.

- *CMPXCHG8B* instruction does not always write to memory (see *QEMU CMPXCHG8B* demo)

  This code works by executing registering a Page Fault handler then executing a *CMPXCHG8B* instruction on a read-only page. For a real CPU, the *CMPXCHG8B* instruction always writes to memory, no matter what is the result. For a read-only page, a Page Fault will be raised. For *QEMU*, no Page Fault occurs. This detection method is available due to what appears to be an oversight.

- Double Fault exception is not supported (see *QEMU* EXC_DBL demo)

  This code begins by setting the limit of the interrupt descriptor table less than what is required to describe the General Protection Fault handler. Then a General Protection Fault is raised. For a real CPU, being unable to raise the General Protection Fault causes the Double Fault exception to be raised. For *QEMU*, the General Protection Fault is raised repeatedly. This detection method is available due to a limitation in the exception handling code.

XII. DETECTING *ATLANTIS* AND *SANDBOX*

Since both *Atlantis* and *Sandbox* emulate only a subset of all of the possible *Windows* APIs, and of those, some of the APIs do not behave in the same way as on a real machine. Thus, they are vulnerable to detection through the use of any unimplemented API or any API that is not emulated correctly. An example is the Beep() API, which has limitations on the frequency of the sound to produce when executed on *Windows NT* and later versions of *Windows*. *Atlantis* does not check that parameter since it emulates *Windows 9x*. Thus, it returns no error, no matter what value is specified. Any program that assumes it is running on *Windows NT* or later will know immediately if *Atlantis* is hosting the environment, by calling that API with an illegal value. Another example is through the use of an exploit. There are several current documented [28] denial-of-service vulnerabilities in different versions of *Windows* for the *Windows* Meta File (WMF) format. If such a malformed WMF file is played successfully, then an operating system emulator is running.

XIII. CONCLUSION

So what can we do? The answer to this question depends on the application that is being used. However, for the reduced privilege guest virtual machines emulators, the ultimate answer is "nothing". The problem for them is that

---

[4] The *REP* instruction is handled specially by x86 CPUs, such that it completes even if the sequence is replaced in memory. For example,

```
mov al, 90h
mov cx, 7
mov di, offset $
rep stosb
jmp $
```

Here, the *NOP* instruction in the AL register is used to overwrite the *REP STOSB* and the following *JMP* instruction. Incorrect emulation (or single-stepping through the code, as with a debugger) will cause the *REP* to exit prematurely, resulting in the *JMP* instruction being executed.

their design does not allow them to intercept non-sensitive instructions that cause information leakage, such as the *SIDT* instruction. As a result, they cannot hide their presence from the RedPill, LDT, and Scooby Doo, attacks.

The Liston/Skoudis paper has a title that suggests that they can reduce the ability of software to detect virtual machine emulators [29]. However, it is actually more concerned with ways to detect virtual machine emulators. The recommendations in that paper for reducing the ability of software to detect virtual machine emulators are exclusively for *VMware*, and insufficient, as noted earlier.

VirtualPC could be improved to intercept the *SIDT* instruction. This would go a long way towards hiding its presence, but it would also need to implement a check for the maximum instruction length.

The interception of the *CPUID* instruction in both *VirtualPC* and *QEMU* to replace the processor identification string should be removed, too.

The use of session key authentication to control guest-to-host and host-to-guest communication in *Parallels* is a good idea that other applications could use.

Bochs, *Hydra*, and *QEMU*, all suffer from bugs and limitations that allow their detection. These are problems that are relatively easily fixed. Given that, only pure software virtual machine emulators can approach complete transparency. It should be possible, at least in theory, to reach the point where detection is unreliable because it can also be attributed to anomalous behavior of a real CPU (for example, the f0 0f bug [30]).

On the other hand, if a majority of future machines run a virtual machine emulator, then malicious code that chooses to not run in its presence will eventually be unintentionally choosing to not run at all.

APPENDIX

VIRTUALPC TIME DEMO:

```
.model  tiny
.code
org     100h

demo:   mov     ax, 2506h
        mov     dx, offset int06
        int     21h
        db      0fh, 3fh, 3, 0
        jmp     $ ;detected
int06:  int     20h
end     demo
```

VIRTUALPC ILEN DEMO:

```
.model  tiny
.code
```

```
org     100h

demo:   mov     ax, 250dh
        mov     dx, offset int0d
        int     21h
        db      0eh dup (2eh)
        jmp     $ ;detected
int0d:  int     20h
end     demo
```

VIRTUALPC BOARD DEMO:

```
For             Each            board           in
GetObject("winmgmts:!\\.\root\cimv2").ExecQuery("Sel
ect * from Win32_BaseBoard")
    If board.Manufacturer = "Microsoft Corporation"
then while 1 : wend 'detected
Next
```

BOCHS WBINVD DEMO:

```
.model  tiny
.486p
.code
org     100h

demo:   mov     edx, ds
        mov     cx, 1000h
        movzx   eax, cx
        add     ah, dh
        mov     es, ax
        shl     eax, 4
        mov     cr3, eax
        shl     edx, 4
        mov     bx, offset gdt
        add     [bx + 2], edx
        mov     [bx + 0ah], dx
        add     [bx+offset idtr-offset gdt+2],edx
        bswap   edx
        mov     [bx + 0ch], dh
        mov     [bx + 0fh], dl
        add     eax, 1007h
        xor     di, di
        stosd
        push    7
        pop     eax
        mov     di, cx
create_tbl:
        stosd
        add     eax, 1000h
        loop    create_tbl
        mov     fs, cx
        cli
        sidt    fword ptr [offset idt_end]
        lidt    [bx + offset idtr - offset gdt]
        lgdt    [bx]
        mov     eax, cr0
        mov     ecx, eax
        or      eax, 80000001h
        mov     cr0, eax
        int     3
int03:  mov     al, fs:[1000h]
        dec     byte ptr es:[1004h]
        wbinvd
        mov     al, fs:[1000h]
        mov     cr0, ecx
        lidt    fword ptr [offset idt_end]
        mov     ah, 4ch
        int     21h
int0e:  jmp     $ ;detected
```

```
gdt     dw      offset gdt_end - offset gdt - 1
        dw      offset gdt
        dd      0
        dd      0ffffh
        dd      9b00h
gdt_end:

idtr    dw      offset idt_end - offset idt - 1
        dw      offset idt
        dw      0

idt     dd      6 dup (0)
        dw      offset int03
        dd      86000008h
        dd      14h dup (0)
        dw      offset int0e
        dd      86000008h
        dw      0
idt_end:

end     demo
```

BOCHS CMPS DEMO:

```
.model  tiny
.code
org     100h

demo:   mov     ax, 2501h
        mov     dx, offset int01
        int     21h
        mov     cx, 101h
        mov     si, cx
        mov     di, cx
        push    cx
        popf
        repe    cmpsb
int01:  jnb     $ ;detected
        int     20h
end     demo
```

BOCHS SCAS DEMO:

```
.model  tiny
.code
org     100h

demo:   mov     ax, 2501h
        mov     dx, offset int01
        int     21h
        mov     cx, 101h
        mov     di, cx
        push    cx
        popf
        repe    scasb
int01:  jnb     $ ;detected
        int     20h
end     demo
```

BOCHS CPUID_AMD1 DEMO:

```
.model  tiny
.586
.code
org     100h

demo:   xor     eax, eax
        cpuid
        cmp     ecx, 444d4163h
```

```
        jne     exit
        mov     eax, 80000000h
        cpuid
        cmp     eax, 2
        jb      exit
        mov     eax, 80000002h
        cpuid
        shr     edx, 1eh
        jb      $ ;detected
exit:   ret
end     demo
```

BOCHS and *QEMU CPUID_AMD2* DEMO:

```
.model  tiny
.586
.code
org     100h

demo:   xor     eax, eax
        cpuid
        cmp     ecx, 444d4163h
        jne     exit
        mov     eax, 8fffffffh
        cpuid
        jecxz   $ ;detected
exit:   ret
end     demo
```

BOCHS ARPL DEMO:

```
.model  tiny
.486p
.code
org     100h

demo:   mov     eax, ds
        shl     eax, 4
        mov     bx, offset gdt
        add     [bx + 2], eax
        mov     [bx + 0ah], ax
        bswap   eax
        mov     [bx + 0ch], ah
        mov     [bx + 0fh], al
        cli
        lgdt    [bx]
        mov     eax, cr0
        inc     ax
        mov     cr0, eax
        cdq
        push    cs
        push    dx
        push    8
        push    offset pmode
        retf
pmode   db      66h
        arpl    dx, ax
        test    edx, edx
        js      $ ;detected
        dec     ax
        mov     cr0, eax
        retf

gdt     dw      offset gdt_e - offset gdt - 1
        dw      offset gdt
        dd      0
        dd      0ffffh
        dd      9b00h
        dd      0ffffh
        dd      0cf9300h
gdt_e:
```

```
end     demo
```

BOCHS and *HYDRA* SEGLOAD DEMO:

```
.model  tiny
.code
org     100h

demo:   mov     ax, 250dh
        mov     dx, offset int0d
        int     21h
        lds     ax, ds:[0fffeh]
        ret
int0d:  jmp     $ ;detected
end     demo
```

HYDRA MOVS DEMO:

```
.model  tiny
.486p
.code
org     100h

demo:   mov     edx, ds
        mov     cx, 1000h
        movzx   eax, cx
        add     ah, dh
        mov     es, ax
        shl     eax, 4
        mov     cr3, eax
        shl     edx, 4
        mov     bx, offset gdt
        add     [bx + 2], edx
        mov     [bx + 0ah], dx
        add     [bx+offset idtr-offset gdt+2],edx
        bswap   edx
        mov     [bx + 0ch], dh
        mov     [bx + 0fh], dl
        add     eax, 1007h
        xor     di, di
        stosd
        push    7
        pop     eax
        mov     di, cx
create_tbl:
        stosd
        add     eax, 1000h
        loop    create_tbl
        mov     fs, cx
        cli
        sidt    fword ptr [offset idt_end]
        lidt    [bx + offset idtr - offset gdt]
        lgdt    [bx]
        mov     eax, cr0
        mov     edx, eax
        mov     ecx, 80000001h
        or      eax, ecx
        mov     cr0, eax
        int     3
int03:  dec     byte ptr es:[1004h]
        xor     esi, esi
        db      64h
        db      67h
        rep     movsw ;shut down Hydra
int0e:  mov     cr0, edx
        lidt    fword ptr [offset idt_end]
        mov     ah, 4ch
        int     21h

gdt     dw      offset gdt_end - offset gdt - 1
        dw      offset gdt
        dd      0
```

```
        dd      0ffffh
        dd      9b00h
gdt_end:

idtr    dw      offset idt_end - offset idt - 1
        dw      offset idt
        dw      0

idt     dd      6 dup (0)
        dw      offset int03
        dd      86000008h
        dw      0
        dd      14h dup (0)
        dw      offset int0e
        dd      86000008h
        dw      0
idt_end:

end     demo
```

HYDRA STOS DEMO:

```
.model  tiny
.486p
.code
org     100h

demo:   mov     edx, ds
        mov     cx, 1000h
        movzx   eax, cx
        add     ah, dh
        movzx   esi, ah
        mov     es, ax
        shl     eax, 4
        mov     cr3, eax
        shl     edx, 4
        mov     bx, offset gdt
        add     [bx + 2], edx
        mov     [bx + 0ah], dx
        add     [bx+offset idtr-offset gdt+2],edx
        bswap   edx
        mov     [bx + 0ch], dh
        mov     [bx + 0fh], dl
        add     eax, 1007h
        xor     di, di
        stosd
        push    7
        pop     eax
        mov     di, cx
create_tbl:
        stosd
        add     eax, 1000h
        loop    create_tbl
        cli
        sidt    fword ptr [offset idt_end]
        lidt    [bx + offset idtr - offset gdt]
        lgdt    [bx]
        mov     eax, cr0
        mov     edx, eax
        mov     ecx, 80000001h
        or      eax, ecx
        mov     cr0, eax
        int     3
int03:  dec     byte ptr es:[esi*4 + 1018h]
        db      67h
        rep     stosw ;shut down Hydra
int0e:  mov     cr0, edx
        lidt    fword ptr [offset idt_end]
        mov     ah, 4ch
        int     21h

gdt     dw      offset gdt_end - offset gdt - 1
        dw      offset gdt
        dd      0
```

```
        dd      0ffffh
        dd      9b00h
gdt_end:

idtr    dw      offset idt_end - offset idt - 1
        dw      offset idt
        dw      0

idt     dd      6 dup (0)
        dw      offset int03
        dd      86000008h
        dw      0
        dd      14h dup (0)
        dw      offset int0e
        dd      86000008h
        dw      0
idt_end:

end     demo
```

QEMU CPUID_AMD DEMO:

```
.model  tiny
.586
.code
org     100h

demo:   xor     eax, eax
        cpuid
        cmp     ecx, 444d4163h
        jne     exit
        mov     eax, 80000000h
        cpuid
        cmp     eax, 2
        jb      exit
        mov     eax, 80000002h
        cpuid
        cmp     eax, 554d4551h
        je      $ ;detected
exit:   ret
end     demo
```

QEMU CMPXCHG8B DEMO:

```
.model  tiny
.586p
.code
org     100h

demo:   mov     edx, ds
        mov     cx, 1000h
        movzx   eax, cx
        add     ah, dh
        mov     es, ax
        shl     eax, 4
        mov     cr3, eax
        shl     edx, 4
        mov     bx, offset gdt
        add     [bx + 2], edx
        mov     [bx + 0ah], dx
        add     [bx+offset idtr-offset gdt+2],edx
        bswap   edx
        mov     [bx + 0ch], dh
        mov     [bx + 0fh], dl
        add     eax, 1007h
        xor     di, di
        stosd
        push    7
        pop     eax
        mov     di, cx
create_tbl:
```

```
        stosd
        add     eax, 1000h
        loop    create_tbl
        mov     fs, cx
        cli
        sidt    fword ptr [offset idt_end]
        lidt    [bx + offset idtr - offset gdt]
        lgdt    [bx]
        mov     eax, cr0
        mov     ecx, eax
        or      eax, 80010001h
        mov     cr0, eax
        int     3
int03:  mov     byte ptr es:[1004h], 5
        mov     al, fs:[1000h]
        inc     ax
        cmpxchg8b fs:[1000h]
        jmp     $ ;detected
int0e:  mov     cr0, ecx
        lidt    fword ptr [offset idt_end]
        mov     ah, 4ch
        int     21h

gdt     dw      offset gdt_end - offset gdt - 1
        dw      offset gdt
        dd      0
        dd      0ffffh
        dd      9b00h
gdt_end:

idtr    dw      offset idt_end - offset idt - 1
        dw      offset idt
        dw      0

idt     dd      6 dup (0)
        dw      offset int03
        dd      86000008h
        dw      0
        dd      14h dup (0)
        dw      offset int0e
        dd      86000008h
        dw      0
idt_end:

end     demo
```

QEMU EXC_DBL DEMO:

```
.model  tiny
.486p
.code
org     100h

demo:   mov     eax, ds
        shl     eax, 4
        mov     bx, offset gdt
        add     [bx + 2], eax
        mov     [bx + 0ah], ax
        add     [bx+offset idtr-offset gdt+2],eax
        bswap   eax
        mov     [bx + 0ch], ah
        mov     [bx + 0fh], al
        cli
        sidt    fword ptr [offset idt_end]
        lidt    [bx + offset idtr - offset gdt]
        lgdt    [bx]
        mov     eax, cr0
        inc     ax
        mov     cr0, eax
        int     3
int03:  int     0ffh
int08:  dec     ax
        mov     cr0, eax
        lidt    fword ptr [offset idt_end]
```

```
        mov     ah, 4ch
        int     21h

gdt     dw      offset gdt_end - offset gdt - 1
        dw      offset gdt
        dd      0
        dd      0ffffh
        dd      9b00h
gdt_end:

idtr    dw      offset idt_end - offset idt - 1
        dw      offset idt
        dw      0

idt     dd      6 dup (0)
        dw      offset int03
        dd      86000008h
        dw      0
        dd      8 dup (0)
        dw      offset int08
        dd      86000008h
        dw      0
idt_end:

end     demo
```

REFERENCES

[1]   Peter Ferrie
      http://pferrie.tripod.com/#hydra
[2]   Methyl
      "Tunneling with Single step mode"
      http://vx.netlux.org/lib/vme04.html
[3]   Methyl
      "Development of Emulation Systems"
      http://vx.netlux.org/lib/vme01.html
[4]   Microsoft
      http://www.microsoft.com/windows/virtualpc
[5]   Hewlett-Packard Laboratories and MIT
      http://www.cag.lcs.mit.edu/dynamorio
[6]   *VMware*
      http://www.vmware.com
[7]   University of Cambridge
      http://www.cl.cam.ac.uk/Research/SRG/netos/xen
[8]   Parallels
      http://www.parallels.com
[9]   Virtuozzo
      http://www.virtuozzo.com
[10] Bugcheck
      "Detecting hardware assisted hypervisors"
      http://www.rootkit.com/newsread.php?newsid=548
[11] Peter Ferrie
      "Detecting hardware-assisted hypervisors without external timing"
      http://www.symantec.com/enterprise/security_response/weblog/2006/09/
      detecting_hardwareassisted_hyp.html
[12] Microsoft
      http://www.microsoft.com/windowsserversystem/virtualserver
[13] Kevin Lawton et al
      http://bochs.sourceforge.net
[14] Fabrice Bellard
      http://fabrice.bellard.free.fr/qemu
[15] Peter Ferrie
      http://pferrie.tripod.com/#atlantis
[16] Norman
      http://www.norman.com
[17] Samuel T. King, Peter M. Chen, Yi-Min Wang, Chad Verbowski, Helen
      J. Wang, and Jacob R. Lorch
      http://www.eecs.umich.edu/virtual/papers/king06.pdf
[18] Joanna Rutkowska
      "Subvirting Vista Kernel For Fun and Profit"
      http://www.whiteacid.org/misc/bh2006/070_Rutkowska.pdf
[19] Dino A. Dai Zovi
      "Hardware Virtualization Rootkits"
      http://www.whiteacid.org/misc/bh2006/036_Zovi.pdf
[20] Danny Quist and Val Smith
      "Detecting the Presence of Virtual Machines Using the Local Data
      Table"
      http://www.offensivecomputing.net/files/active/0/vm.pdf
[21] Joanna Rutkowska
      "Red Pill"
      http://invisiblethings.org/papers/redpill.html
[22] Tobias Klein
      "Scooby Doo - VMware Fingerprint Suite"
      http://www.trapkit.de/research/vmm/scoopydoo/index.html
[23] Tobias Klein
      "jerry - A(nother) VMware Fingerprinter"
      http://www.trapkit.de/research/vmm/jerry/index.html
[24] Peter Ferrie
      "Tumours and Polips"
      http://pferrie.tripod.com/vb/polip.pdf
[25] Ken Kato et al
      "VMware Backdoor I/O Port"
      http://chitchat.at.infoseek.co.jp/vmware/backdoor.html
[26] Tim Shelton
      http://lists.grok.org.uk/pipermail/full-disclosure/2005-
      December/040442.html
[27] Ben Armstrong
      "Detecting Microsoft virtual machines"
      http://blogs.msdn.com/virtual_pc_guy/archive/2005/10/27/484479.aspx
[28] Peter Ferrie
      "Inside the Windows Meta File Format"
      http://pferrie.tripod.com/vb/wmf.pdf
[29] Tom Liston and Ed Skoudis
      "On the Cutting Edge: Thwarting Virtual Machine Detection"
      http://handlers.sans.org/tliston/ThwartingVMDetection_Liston_Skoudis.
      pdf
[30] Robert R. Collins
      "The Intel Pentium F00F Bug Description and Workarounds"
      http://www.x86.org/errata/dec97/f00fbug.htm